

# Some mathematics of Word2Vec algorithm and approximated Softmax

A/Prof Richard Yi Da Xu, Erica Huang  
richardxu.com



University of Technology Sydney (UTS)

July 15, 2018

# word embedding

- ▶ words are symbols: one may **not** able to perform arithmetic operations on them.
- ▶ turning each word into a **vector**, e.g., “machine” → [2.4 - 1.2 1.9 ...]
- ▶ can measure how similar or dissimilar between them
- ▶ can even perform “arithmetic”:
- ▶ examples from the original paper:

$$\text{vec}(\text{King}) - \text{vec}(\text{man}) + \text{vec}(\text{woman}) = \text{vec}(\text{Queen})$$



*Mikolov et. al., (2013) “Linguistic Regularities in Continuous Space Word Representations”*

- ▶ simple!

$$\begin{bmatrix} \text{"a"} & 1 & 0 & \dots & 0 \\ \text{"abbreviate"} & 0 & 1 & \dots & 0 \\ \vdots & & & & \\ \text{"zoology"} & 0 & 0 & \dots & 1 \end{bmatrix}$$

- ▶ however, the structure is huge and sparse
- ▶ every pair of items are  $\sqrt{2}$  apart.
- ▶ can we do better?

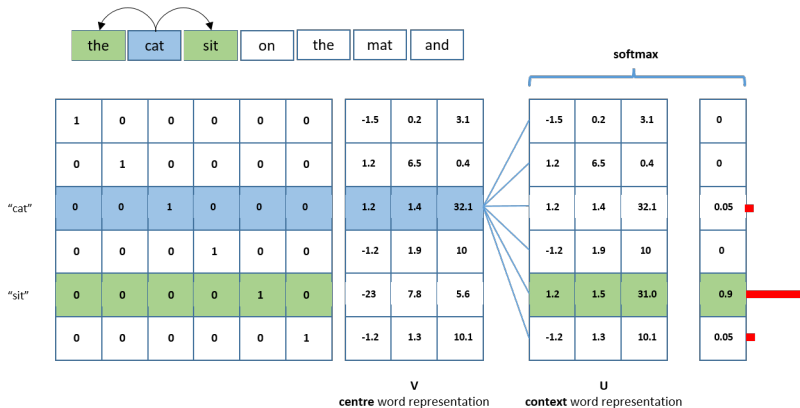
## supervised learning between (target $\rightarrow$ context)

- ▶ Word2Vec algorithm tries to leverage ("*target*", "*context*") relationships:
- ▶ offers two approaches, i.e., to maximize two different conditional densities:
  - ▶ skip-gram:  $\Pr(\text{"context"} \mid \text{"target"})$
  - ▶ continuous bag of words (CBOW):  $\Pr(\text{"target"} \mid \text{"context"})$
- ▶ uses supervised learning techniques: so we need to build ("*input*", "*label*") pairs
  1. pick window size (odd number)
  2. extract all tokens based on this chosen window size
  3. remove middle word in each window; this becomes your target word, rest are context
- ▶ btw, each word is an object class, a huge softmax!

- ▶ for example, **Skip-gram**(window size 3)
- ▶ *"the cat sit on the mat"*
  1. "the", "cat", "sit", target: **cat**
  2. "cat", "sit", "on", target: **sit**
  3. "sit", "on", "the", target: **on**
  4. "on", "the", "mat", target: **the**
- ▶ now we can perform **supervised learning** for (center, context):
  - ▶ ("cat", "the")
  - ▶ ("cat", "sit")
  - ▶ ("sit", "cat")
  - ▶ ("sit", "on")
  - ▶ ("on", "sit")
  - ▶ ("on", "the")
  - ▶ ("the", "on")
  - ▶ ("the", "mat")


# Simple Skip-gram example

- say vocabulary has 6 unique words in total
- "cat (3<sup>th</sup> word)" and "sit (5<sup>th</sup> word)" is a (center, context) pair
- for every word  $w$ , it has 2 representations  $\mathbf{u}_w$  and  $\mathbf{v}_w$ , one for input and one for output
- predict context word given a center word  $\Pr(o = \text{"sit"} | c = \text{"cat"})$



# Simple CBoW example

- to predict center word given multiple context words:



- $\mathbf{v}_t$  is average of input (context) vectors
- the new objective is:

$$p(c|t) = \frac{\exp(\mathbf{u}_c^\top \mathbf{v}_t)}{\sum_{w'} \exp(\mathbf{u}_{w'}^\top \mathbf{v}_t)}$$

# Skip-Gram objective function (1)

- predict context word given a center word  $\Pr(o = \text{"sit"} | c = \text{"cat"})$

$$\Pr(o = \text{"sit"} | c = \text{"cat"}) = \frac{\exp(\mathbf{u}_{\text{"sit"}}^\top \mathbf{v}_{\text{"cat"}})}{\sum_{w \in \mathcal{V}} \exp(\mathbf{u}_w^\top \mathbf{v}_{\text{"cat"}})}$$
$$\Rightarrow \log(\Pr(o = | c)) = \log \left( \frac{\exp(\mathbf{u}_o^\top \mathbf{v}_c)}{\sum_{w \in \mathcal{V}} \exp(\mathbf{u}_w^\top \mathbf{v}_c)} \right)$$

- we need to compute both  $\frac{\partial \log(\Pr(o = | c))}{\partial \mathbf{v}_c}$  and  $\frac{\partial \log(\Pr(o = | c))}{\partial \mathbf{u}_w}$ ,  $\forall \mathbf{v}_c, \mathbf{u}_w \in \mathcal{V}$
- due to symmetry, looking at only one:

$$\begin{aligned} \frac{\partial \log(\Pr(o = | c))}{\partial \mathbf{v}_c} &= \frac{\partial \mathbf{u}_o^\top \mathbf{v}_c}{\partial \mathbf{v}_c} - \frac{\partial \log \left( \sum_{w \in \mathcal{V}} \exp(\mathbf{u}_w^\top \mathbf{v}_c) \right)}{\partial \mathbf{v}_c} \\ &= \mathbf{u}_o - \left( \frac{\partial}{\partial \mathbf{v}_c} \log \left( \sum_{w \in \mathcal{V}} \exp(\mathbf{u}_w^\top \mathbf{v}_c) \right) \right) \\ &= \mathbf{u}_o - \left( \frac{1}{\sum_{w \in \mathcal{V}} \exp(\mathbf{u}_w^\top \mathbf{v}_c)} \frac{\partial}{\partial \mathbf{v}_c} \left( \sum_{w \in \mathcal{V}} \exp(\mathbf{u}_w^\top \mathbf{v}_c) \right) \right) \\ &= \mathbf{u}_o - \left( \frac{1}{\sum_{w \in \mathcal{V}} \exp(\mathbf{u}_w^\top \mathbf{v}_c)} \left( \sum_{w \in \mathcal{V}} \frac{\partial}{\partial \mathbf{v}_c} \exp(\mathbf{u}_w^\top \mathbf{v}_c) \right) \right) \\ &= \mathbf{u}_o - \frac{1}{\sum_{w \in \mathcal{V}} \exp(\mathbf{u}_w^\top \mathbf{v}_c)} \left( \sum_{w \in \mathcal{V}} \mathbf{u}_w \exp(\mathbf{u}_w^\top \mathbf{v}_c) \right) \\ &= \mathbf{u}_o - \frac{\sum_{w \in \mathcal{V}} \mathbf{u}_w \exp(\mathbf{u}_w^\top \mathbf{v}_c)}{\sum_{w \in \mathcal{V}} \exp(\mathbf{u}_w^\top \mathbf{v}_c)} \end{aligned}$$



## Skip-Gram objective function (2)

- derivative:

$$\begin{aligned}\frac{\partial \log(\Pr(o = |c))}{\partial \mathbf{v}_c} &= \mathbf{u}_o - \frac{\sum_{w \in \mathcal{V}} \mathbf{u}_w \exp(\mathbf{u}_w^\top \mathbf{v}_c)}{\sum_{w \in \mathcal{V}} \exp(\mathbf{u}_w^\top \mathbf{v}_c)} \\&= \mathbf{u}_o - \sum_{w \in \mathcal{V}} \frac{\exp(\mathbf{u}_w^\top \mathbf{v}_c)}{\sum_{w \in \mathcal{V}} \exp(\mathbf{u}_w^\top \mathbf{v}_c)} \mathbf{u}_w \\&= \mathbf{u}_o - \sum_{w \in \mathcal{V}} \Pr(w|c) \mathbf{u}_w \\&= \mathbf{u}_o - \mathbb{E}_{w \sim \Pr(w|c)} [\mathbf{u}_w]\end{aligned}$$

- obviously, there are  $|\mathcal{V}|$  is too big, making it too computationally expensive to compute the sum
- can we do better?

# Negative sampling (1)

- ▶ negative sampling based on Skip-Gram model, it is optimizing **different objective**, let  $\theta = [\mathbf{u}, \mathbf{v}]$ :
- ▶ we let  $\bar{w}$  to indicate **negative samples**, and come from a negative population  $\bar{D}$

$$\begin{aligned}\theta &= \arg \max_{\theta} \prod_{(w,c) \in D} \Pr(D = 1 | w, c, \theta) \prod_{(\bar{w}, c) \in \bar{D}} \Pr(D = 0 | \bar{w}, c, \theta) \\&= \arg \max_{\theta} \prod_{(w,c) \in D} \Pr(D = 1 | w, c, \theta) \prod_{(\bar{w}, c) \in \bar{D}} (1 - \Pr(D = 1 | \bar{w}, c, \theta)) \\&= \arg \max_{\theta} \sum_{(w,c) \in D} \log (\Pr(D = 1 | w, c, \theta)) + \sum_{(\bar{w}, c) \in \bar{D}} \log (1 - \Pr(D = 1 | \bar{w}, c, \theta)) \\&= \arg \max_{\theta} \sum_{(w,c) \in D} \log \frac{1}{1 + \exp [-\mathbf{u}_w^\top \mathbf{v}_c]} + \sum_{(\bar{w}, c) \in \bar{D}} \log \left( 1 - \frac{1}{1 + \exp [-\mathbf{u}_{\bar{w}}^\top \mathbf{v}_c]} \right) \\&= \arg \max_{\theta} \sum_{(w,c) \in D} \sigma(-\mathbf{u}_w^\top \mathbf{v}_c) + \sum_{(\bar{w}, c) \in \bar{D}} \log \left( \frac{1}{1 + \exp [\mathbf{u}_{\bar{w}}^\top \mathbf{v}_c]} \right) \\&= \arg \max_{\theta} \sum_{(w,c) \in D} \sigma(\mathbf{u}_w^\top \mathbf{v}_c) + \sum_{(\bar{w}, c) \in \bar{D}} \log \sigma(-\mathbf{u}_{\bar{w}}^\top \mathbf{v}_c)\end{aligned}$$

# Negative sampling (2)

- ▶ negative sampling based on Skip-Gram model, it is optimizing **different objective**, let  $\theta = [\mathbf{u}, \mathbf{v}]$ :

$$\theta = \arg \max_{\theta} \sum_{(w,c) \in D} \sigma(\mathbf{u}_w^\top \mathbf{v}_c) + \sum_{(\bar{w},c) \in \bar{D}} \log \sigma(-\mathbf{u}_{\bar{w}}^\top \mathbf{v}_c)$$

- ▶ it still has a huge sum term  $\sum_{(\bar{w},c) \in \bar{D}} (\cdot)$ , so we change to:

$$\theta = \arg \max_{\theta} \sigma(\mathbf{u}_w^\top \mathbf{v}_c) + \sum_{\bar{w}=1}^k \mathbb{E}_{\bar{w} \sim P(w)} \log \sigma(-\mathbf{u}_{\bar{w}}^\top \mathbf{v}_c)$$

- ▶ sample a fraction of negative samples in second terms:  $\{\bar{w}\}$  instead of going for  $\forall (\bar{w} \neq w) \in \mathcal{V}$
- ▶  $\bar{w} \sim \text{Pr}_{\bar{D}}(w)$ , where  $\text{Pr}_{\bar{D}}(\cdot)$  is probability of negative sample:
  - can use Unigram Model raised to the power of  $\frac{3}{4}$
- ▶ doing so, we can:
  - increase probability of popular words marginally
  - increase probability of rarer words dramatically
  - making “rarer” words also have chance to be sampled
- ▶ in unigram model, probability of each word only depends on that word's own probability

# Noise Contrastive Estimation (NCE) (1)

- ▶ let  $u_\theta(w, c)$  be un-normalized score function, i.e.,  $u_\theta(w, c) = \exp(\mathbf{u}_w^\top \mathbf{v}_c)$

$$P_\theta(w|c) = \frac{u_\theta(w, c)}{\sum_{w' \in \mathcal{V}} u_\theta(w', c)} = \frac{u_\theta(w, c)}{Z_c}$$

- ▶  $\tilde{p}(w|c)$  and  $\tilde{p}(c)$  are empirical distributions  
we **know** them from data, so we can sample  $(w, c)$  from it!
- ▶ a “noise” distribution  $q(w)$  is used - uniform or uniform unigram  
we also **know** them, again, we can sample  $\tilde{w} \sim q(\cdot)$
- ▶ **task** is to use sample from both distributions, then to assist us find  $\theta$  making  $P_\theta(w|c)$  to approximate empirical distribution as closely as possible (by minimal cross entropy)

# Noise Contrastive Estimation (NCE) (2)

- ▶ **training data generation:**  $(w, c, D) \sim \mathcal{D}$
- ▶ of course, to utilize  $\tilde{p}(w|c)$ ,  $\tilde{p}(c)$  and  $q(w)$ , which we already have knowledge of:
  1. sample a  $c \sim \tilde{p}(c)$ ,  $w \sim \tilde{p}(w|c)$  and label it as  $D = 1$
  2.  $k$  “noise” samples from  $q(\cdot)$ , and label it as  $D = 0$
- ▶ NCE transforms:
  - “problem of model estimation” (computationally expensive) to
  - “problem of estimating parameters of *probabilistic binary classifier* uses **same parameters** to distinguish between samples”: (computationally acceptable)
- ▶ from empirical distribution
- ▶ from noise distribution

# Noise Contrastive Estimation (NCE) (3)

- ▶ let  $u_\theta(w, c)$  be un-normalized score function, i.e.,  $u_\theta(w, c) = \exp(\mathbf{u}_w^\top \mathbf{v}_c)$

$$\begin{aligned} P(D=0|c, w) &= \frac{P(D=0, w|c)}{P(w|c)} = \frac{p(w|D=0, c)P(D=0)}{\sum_{d \in \{0,1\}} p(w|D=d, c)P(D=d)} \\ &= \frac{q(w) \times \frac{k}{1+k}}{\tilde{P}(w|c) \times \frac{1}{k+1} + q(w) \times \frac{k}{1+k}} \\ &= \frac{kq(w)}{\tilde{P}(w|c) + kq(w)} \\ P(D=1|c, w) &= 1 - P(D=0|c, w) \\ &= \frac{\tilde{P}(w|c)}{\tilde{P}(w|c) + kq(w)} \end{aligned}$$

# Noise Contrastive Estimation (NCE) (4)

- NCE replaces empirical distribution  $\tilde{p}(w|c)$  with model distribution  $p_\theta(w|c)$

$$P(D = 0|c, w) = \frac{kq(w)}{\tilde{P}(w|c) + kq(w)} = \frac{kq(w)}{\frac{u_\theta(w|c)}{Z_c} + kq(w)}$$

$$P(D = 1|c, w) = \frac{\tilde{P}(w|c)}{\tilde{P}(w|c) + kq(w)} = \frac{\frac{u_\theta(w|c)}{Z_c}}{\frac{u_\theta(w|c)}{Z_c} + kq(w)}$$

- $\theta$  is then chosen to maximize likelihood of proxy corpus created from **training data generation**:

$$\mathcal{L}^{\text{NCE}} = \log p(D = 1|c, w) + k \sum_{(w,c) \in \mathcal{D}} \mathbb{E}_{w' \sim q} \log p(D = 0|c, w')$$

- for neural networks:  $Z_c$  can also be trained or set to some fixed number, e.g.,  $Z_c = 1$
- **negative sampling** is its special case  $k = |\mathcal{V}|$  and  $q(\cdot)$  is uniform, and  $Z_c = 1$ :

$$P(D = 0|c, w) = \frac{|\mathcal{V}| \frac{1}{|\mathcal{V}|}}{u_\theta(w|c) + |\mathcal{V}| \frac{1}{|\mathcal{V}|}} = \frac{1}{u_\theta(w|c) + 1}$$

$$P(D = 1|c, w) = \frac{u_\theta(w|c)}{u_\theta(w|c) + |\mathcal{V}| \frac{1}{|\mathcal{V}|}} = \frac{u_\theta(w|c)}{u_\theta(w|c) + 1}$$

# Self Normalization

- ▶ in previous slide, we want to normalize, s.t.,  $Z_c = 1$
- ▶ start with  $u_\theta(w, c) = \exp(\mathbf{u}_w^\top \mathbf{v}_c)$ :

$$\begin{aligned} P_\theta(w|c) &= \prod_w \frac{\exp(\mathbf{u}_w^\top \mathbf{v}_c)}{Z_c} \\ \implies J_\theta &= - \prod_w \log(P_\theta(w|c)) = - \sum_w \log\left(\frac{\exp(\mathbf{u}_w^\top \mathbf{v}_c)}{Z_c}\right) \\ &= - \sum_w \mathbf{u}_w^\top \mathbf{v}_c - \log(Z_c) \end{aligned}$$

- ▶ to constrain model and sets  $Z(c) = 1 \implies \log Z(c) = 0$ :

$$\begin{aligned} J_\theta &= - \sum_w \mathbf{u}_w^\top \mathbf{v}_c + \log Z(c) - \alpha (\log(Z(c)) - 0)^2 \\ &= - \sum_w \mathbf{u}_w^\top \mathbf{v}_c + \log Z(c) - \alpha \log^2 Z(c) \end{aligned}$$



A library created by Facebook research team for

1. efficient learning of word representations(Enriching Word Vectors with Subword Information)
2. sentence classification(Bag of Tricks for Efficient Text Classification)

- ▶ So how is it different from Word2Vec?
- ▶ Instead of words, we now have **ngrams** of subwords, what is its **advantage**?
  1. Helpful for finding representations for rare words
  2. Give vector representations for words not present in dictionary
- ▶ for example,  $n = 3$ , i.e., 3-grams:
  - ▶ **word**: “where”,
  - ▶ **sub-words**: “wh”, “whe”, “her”, “ere”, “re”
- ▶ we then represent a word by the **sum of the vector representations** of all its  $n$ -grams
- ▶ to compute an un-normalised score with center word  $\mathbf{v}_c$ , given a word  $w$ ,  $g_w$  is the set of  $n$ -grams appearing in  $w$ ,  $z_g$  is the representation to each individual  $n$ -gram

$$u(w, c) = \exp \left[ \sum_{g \in g_w} z_g^\top \mathbf{v}_c \right]$$

# Global Vectors for Word Representation(GloVe)

- ▶ **co-occurrence probabilities** are useful
- ▶ GloVe learns word vectors through *word co-occurrences*
- ▶ co-occurrence matrix  $P$  where  $P_{ij}$  is how often word  $i$  appears in the context of word  $j$
- ▶ Fast training and scalable to huge corpora
- ▶ loss function:

$$\theta^* = \arg \min_{\theta} \left( J(\theta) \equiv \frac{1}{2} \sum_{\mathbf{u}, \mathbf{v}_j \in \mathcal{V}} f(P_{ij})(\mathbf{u}_i^{\top} \mathbf{v}_j - \log P_{ij})^2 \right)$$


- ▶ it tries to minimize difference:

$$(\mathbf{u}_i^{\top} \mathbf{v}_j - \log P_{ij})$$

- ▶ more frequently two words appear together, more similar their vector representation should be
- ▶  $f(\cdot)$  is weighting function to “prevent” certain scenarios, for example:

$$P_{ij} = 0 \implies \log P_{ij} = -\infty \implies f(0) = 0$$

# Hierarchical Softmax (1)



- ▶ each word  $w_i$  has a unique (pre-defined) path (not a random path!), which performs a left or right turn from nodes:  $n(w_i, 1)$ ,  $n(w_i, 2)$ ,  $n(w_i, 3)$ , . . .
- ▶ the route is defined in such a way that, each child node is from a (LEFT/RIGHT) "channel" of its parent: i.e.,  $n(w, j + 1) = \text{ch}(n(w, j))$
- ▶ for example:


1.  $n(w_2, 2) = \text{LEFT}(n(w_2, 1))$
  2.  $n(w_2, 3) = \text{LEFT}(n(w_2, 2))$
  3.  $n(w_2, 4) = \text{RIGHT}(n(w_2, 3))$
- $\underbrace{\hspace{10em}}_{w_2}$

- ▶ there are  $V$  words in leaf (white node)
- ▶ there are  $V - 1$  inner (non-leaf) nodes (grey node) each associate with a value of  $\mathbf{v}$  which is shared among all words going through this node

Xin Rong, word2vec Parameter Learning Explained

- ▶ **super advantage:**  $\Pr(w|c)$  is already a probability by multiplying all probabilities of path, no need to normalize!

# Hierarchical Softmax (2)



Xin Rong, word2vec Parameter Learning Explained

- **super advantage:**  $\Pr(w|c)$  is already a probability by multiplying all probabilities of path, no need to normalize!

we define:  $\xi[\cdot] = \begin{cases} 1 : & \text{true} \\ -1 : & \text{false} \end{cases}$

$$\Pr(w|c) = \prod_{j=1}^{L(w)-1} \sigma \left( \underbrace{\xi[n(w, j+1) = \text{ch}(n(w, j))]}_{\text{control its sign}} \mathbf{v}_{n(w, j)}^\top \mathbf{u}_c \right)$$

- looking at  $\Pr(w_2|c)$  and  $\Pr(w_3|c)$ :
- $n(w_2, 1) = n(w_3, 1)$  in fact  $\{n(w_i, 1)\}_{i=1}^{|\mathcal{V}|}$  all equal
- $n(w_2, 2) = n(w_3, 2)$

$$\begin{aligned} \Pr(w_2|c) &= p(n(w_2, 1), \text{LEFT}) p(n(w_2, 2), \text{LEFT}) p(n(w_2, 3), \text{RIGHT}) \\ &= \sigma \left( \mathbf{v}_{n(w_2, 1)}^\top \mathbf{u}_c \right) \sigma \left( \mathbf{v}_{n(w_2, 2)}^\top \mathbf{u}_c \right) \sigma \left( -\mathbf{v}_{n(w_2, 3)}^\top \mathbf{u}_c \right) \\ \Pr(w_3|c) &= p(n(w_3, 1), \text{LEFT}) p(n(w_3, 2), \text{RIGHT}) \\ &= \sigma \left( \mathbf{v}_{n(w_3, 1)}^\top \mathbf{u}_c \right) \sigma \left( -\mathbf{v}_{n(w_3, 2)}^\top \mathbf{u}_c \right) \end{aligned}$$

- **consideration 1**  $\exp(\mathbf{x}^T \boldsymbol{\theta}_j)$  can become very large:

$$\begin{aligned}\pi_j &= \frac{\exp(\mathbf{x}^T \boldsymbol{\theta}_j)}{\sum_{l=1}^3 \exp(\mathbf{x}^T \boldsymbol{\theta}_l)} \\&= \frac{(\exp(\mathbf{x}^T \boldsymbol{\theta}_j)) / C}{(\sum_{l=1}^3 \exp(\mathbf{x}^T \boldsymbol{\theta}_l)) / C} = \frac{\exp(\mathbf{x}^T \boldsymbol{\theta}_j - C)}{\sum_{l=1}^3 \exp(\mathbf{x}^T \boldsymbol{\theta}_l - C)} \\&= \frac{\exp(\mathbf{x}^T \boldsymbol{\theta}_j - \max(\{\exp(\mathbf{x}^T \boldsymbol{\theta}_l)\}))}{\sum_{l=1}^3 \exp(\mathbf{x}^T \boldsymbol{\theta}_l - \max(\{\exp(\mathbf{x}^T \boldsymbol{\theta}_l)\}))}\end{aligned}$$

- **consideration 2** arg max operation, can be done without exp, i.e.,

$$\arg \max_{i \in \{1, \dots, k\}} (\pi_1, \dots, \pi_k) \equiv \arg \max_{i \in \{1, \dots, k\}} (\mathbf{x}^T \boldsymbol{\theta}_1, \dots, \mathbf{x}^T \boldsymbol{\theta}_k)$$